

# NGINX CORE

*Flawless Application Delivery*

# Prerequisites/Expectations

- Sysadmin, DevOps, Solution Architect
- Some familiarity with Web Servers
- Some familiarity with Linux
- Text Editor: Vim, Vi, Emacs etc.
- Some knowledge of Networking

# The Training Environment

- AWS EC2 Instances
- Ubuntu
- NGINX Plus

# Log Into AWS

If you haven't done so already, please take the time to SSH into your EC2 Instances (Windows users use **PuTTY**).

Check your email for the login credentials, check your spam folder!

```
ssh student<number>@<ec2-server-hostname>
```

# Course Administration

- Course Duration: 7-8 hours
- 10 minute break at the top of each hour
- 30-60 min break after part 1
- Ask questions at any time!

# Agenda: Part One



- Overview
- Serving Static Content
- Proxying Connections
- Logging
- Security
- Variables

# NGINX OVERVIEW

# Module Objectives

This module enables you to:

- Gain a basic understanding of NGINX's features
- Learn about the history of NGINX
- Understand the various use cases that NGINX supports




# Origins of NGINX




- 2002: Igor Syrov working for rambler.ru
- 2004: First OSS release
- 2011: Company founded
- 2016-Present: 500+ customers and 80+ employees


# High Concurrency



# Low Memory



# NGINX Use Cases



HTTP traffic



Reverse Proxy

*Caching, Load Balancing...*



Webserver

*Serve content from disk*

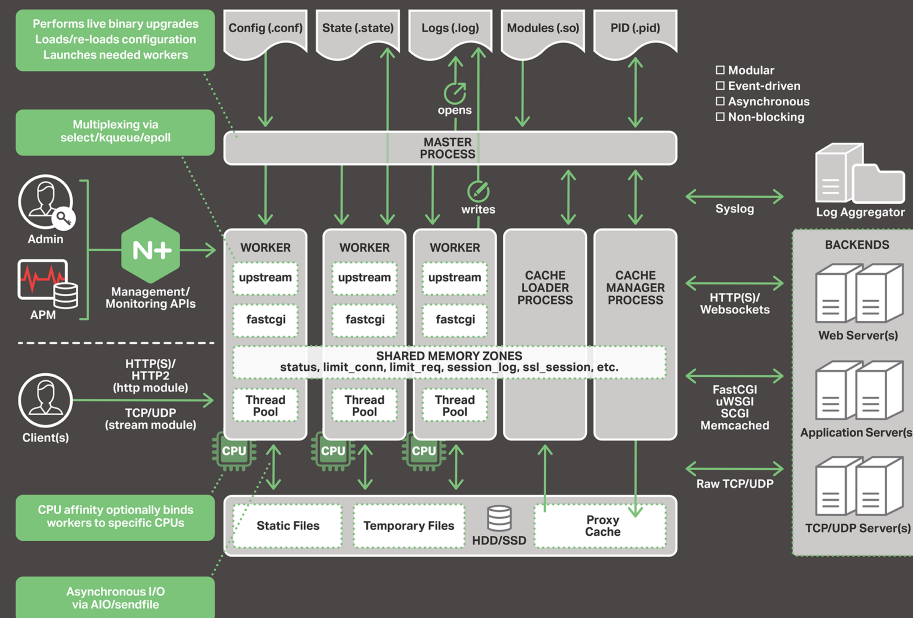


Application Gateway

*FastCGI, uWSGI, Passenger...*

# NGINX Architecture

## The Powerful and Efficient Architecture of NGINX



**Master:** Evaluates config, modules, PID, logs

**Workers:** Handles requests and responses

**Shared Memory:** Counters, rate limits, status, session log etc.

# Basic NGINX Commands

#reloads config

\$ nginx -s reload

#graceful shutdown

\$ nginx -s quit

#terminates nginx process

\$ nginx -s stop

#config syntax check (pre-reload)

\$ nginx -t

#displays currently running configs


\$ nginx -T

#checks nginx version

\$ nginx -v

# Reloading the Configuration

1. `nginx -s reload` sends `SIGHUP` signal to kernel
2. Master evaluates new config
3. Checks for `emerg` level errors
4. Forks new workers while old workers gracefully shut down



# SERVING STATIC CONTENT



# Module Objectives

This module enables you to:

- Understand the Configuration File
- Configure a Basic Setup
- Explore server selection methods

# Running Processes

To check running processes, run the following command:

```
$ ps aux | grep nginx
```

# Path to Files

## Executable Path

```
$ /usr/sbin/nginx
```

## Log Path

```
$ /var/log/nginx
```

# Configuration File

Global Configuration Path

```
$ /etc/nginx/nginx.conf
```

Additional Configuration(s) Path

```
$ /etc/nginx/conf.d/*.conf
```

Documentation: [nginx.conf Example](#)

# Include Directive

The following line in `nginx.conf` allows NGINX to search for additional configurations

```
include /etc/nginx/conf.d/*.conf;
```

# Configuration File Structure

- Directives
- Blocks
- Contexts

Documentation: [How .conf Files Work](#)

# Directives

Configuration statement that controls **NGINX Modules**

```
listen 80;  
root /usr/share/nginx/html;  
index index.html index.htm index.php;
```

# Blocks

Contains mixture of directives and data—begins and ends with curly brackets.

```
server {  
    listen 80;  
    root /usr/share/nginx/html;  
    index index.html index.htm index.php;  
}
```



# Contexts

Nested Blocks implying a hierarchy. Colloquially, 'Block' and 'Context' are interchangeable.

```
http {  
    include      /etc/nginx/mime.types;  
    default_type application/octet-stream;  
    gzip on;  
    log_format  
  
    server {  
        listen 80;  
        root /usr/share/nginx/html;  
        index index.html index.htm index.php;  
  
        location {  
            proxy_pass http://backend;  
        }  
    }  
}
```

Documentation: [Further Explanation](#)

# Serving Content

Requirements:

- `http` - high level processing (logging, compression, caching etc.)
- `server` - virtual server that handles the request
- `location` - processing based on request URI

# **server** Block Example

- Defines virtual **server** ("VirtualHost" in Apache)
- Always nested inside either **http** or **stream** context
- Binds to TCP sockets with **server\_name** and **listen**

```
server {  
    listen 80;  
    server_name www.example.com;  
    root /etc/student1/public_html;  
}
```

Documentation: [Server Block Examples](#)

# **listen** Directive

- Defines IP / Port that **server** responds to
- Default is **0.0.0.0:80** (**:8080** for non-root)
- Can be: IP, IP:Port, Port, Unix Socket

# listen Rules

NGINX will substitute values for "incomplete" directives

- No `listen` directive uses `0.0.0.0:80`
- IP with no port uses `<some_ip>:80`
- Port with no IP uses `0.0.0.0:<some_port>`

IP takes priority over port

# listen Example

If example.com is hosted on port 80 of 192.168.1.10, the first block serves the response

```
server {  
    listen 192.168.1.10;  
  
}  
  
server {  
    listen 80;  
    server_name example.com;  
  
}
```

# **server\_name** Directive

Used to differentiate between matching IP:Port specificity

- Checks the "Host" header field
- Matching Rules (in order):
  - Exact Match
  - Leading Wildcard
  - Trailing Wildcard
  - Regex

# **server\_name** Example 1

If "Host" value matches "host1.example.com" exactly, second block serves response

```
server {  
    listen 80;  
    server_name *.example.com;  
  
}  
  
server {  
    listen 80;  
    server_name host1.example.com;  
  
}
```



## server\_name Example 2

If "Host" value is "www.example.org", second block serves response — if no match, then trailing wildcards take over

```
server {  
    listen 80;  
    server_name www.example.*;  
  
}
```

```
server {  
    listen 80;  
    server_name *.example.org;  
  
}
```

```
server {  
    listen 80;  
    server_name *.org;  
  
}
```

## **server\_name** Example 3

The FIRST matching regex has priority. If the request is "www.example.com", the first block serves the response

```
server {  
    listen 80;  
    server_name ~^(www|host1).*\.example\.com$;  
  
}
```

```
server {  
    listen 80;  
    server_name ~^(subdomain|set|www|host1).*\.example\.com$;  
  
}
```

# IP vs. **server\_name** Recap

If NGINX cannot determine which server to select, it will choose the first block in the configuration

```
server {  
    listen 192.168.1.1:80;  
    server_name www.example.*;  
}  
  
server {  
    listen 192.168.1.1:80;  
    server_name *.example.net;  
}  
  
server {  
    listen 192.168.1.1:80;  
    server_name ~^(www|host1).*\.example\.com$;  
}
```

# default\_server

- Checks against "HOST" header field
- Overrides "first-found" algorithm
- Only one declaration per IP:Port combo

```
server {  
    listen 80 default_server;  
    server_name example.net www.example.net;  
}
```

# Bad Requests

`server_name` with an empty string value can prevent bad requests. 444 response code closes TCP connection

```
server {  
    listen 80;  
    server_name "";  
    return 444;  
}
```

# Location Block Example

Breaks down request further by URI

```
location /application1 {  
  
}
```

Two most common types:

- **Prefix**
- **Regex**

# Prefix Location

- Checked first, then longest match serves response
- Nested inside **server** context

```
location /application1 {  
  
}  
  
location /application1/images {  
    alias /media/data;  
}
```

Second prefix serves response if request is:

```
$ curl http://somedomain.com/application1/images/?img2
```

# Location Modifiers

- Exact String Match =
- Case Insensitive Regex ~\*
- Case Sensitive Regex ~
- Stop Request Processing ^~
- Named Location Routing @



# Regex Location

Matched sequentially and only after prefix locations.

```
location /application1 {  
  
}  
  
location ~* ^\.(gif|jpg|jpeg|png)$ {  
    alias /media/data;  
}
```

# Location Order

## Configuration Example

```
server {  
    listen 80 default_server;  
    root /usr/share/nginx/html;  
  
    location = / {  
    }  
  
    location ~* ^\.(png|jpg)$ {  
    }  
  
    location ^~ /app1 {  
    }  
}
```

## Selection Order:

- Location 1
- Location 3
- Location 2

# Lab 1.1: Serve Pages and Images

Note: commands in "/etc/nginx/conf.d" require "sudo" or "root" level privileges

## 1. Navigate to the NGINX configuration directory:

```
$ cd /etc/nginx/conf.d
```

## 2. Backup the default and ssl configurations:

```
$ sudo mv default.conf default.conf.bak  
$ sudo mv example_ssl.conf example_ssl.conf.bak
```

## 3. Create a new configuration called `server1.conf`:

```
$ sudo vim server1.conf
```

# Lab 1.2: Serve Pages and Images

1. Create the following `server` block:

```
server {  
    listen 80;  
    root /home/student1/public_html;  
}
```

2. Add the following three `location` prefixes:

```
location /application1 {}  
location /application2 {}  
location /images {}
```

3. Add an overriding `root` directive in `/images`:

```
root /data;
```

# PROXYING CONNECTIONS

# Module Objectives

This module enables you to:

- Configure Proxy Server
- Understand how Proxy Buffering works
- Demonstrate use of `proxy_pass` directive

# Reverse Proxy Servers

Receives requests, passes them to backend servers

NGINX supports proxy for HTTP, HTTPS, TCP, UDP, and other protocols.

```
server {  
    listen 80;  
    server_name mydomain.com;  
  
    location / {  
        proxy_pass http://backend;  
    }  
}
```

# proxy\_pass Directive

Sets the address, and protocol, of the proxied server(s)

Syntax:

```
proxy_pass $protocol$host$uri
```

Example:

```
location / {  
    proxy_pass http://backend:8080/application/  
}
```



# Proxy With a Path

Replaces matched portion of the request uri, including location prefix:

```
#curl request example:
$ curl http://server/application1/

#NGINX selects a matching location prefix
location /application1 {
    proxy_pass http://localhost:8080/otherapp;
}

#proxy request becomes
http://server:8080/otherapp/app1.html;
```

# Proxy Without a Path

Replaces matched portion of the request uri, NOT including location prefix:

```
#curl request example:
$ curl http://server/application2/

#NGINX selects a matching location prefix
location /application2 {
    proxy_pass http://localhost:8080
}

#proxy request becomes
http://server:8080/application2/index.html
```

# To / or Not To /

Beware of trailing slash request mismatches.

```
#Example curl request with NO trailing slash
$ curl http://example.com/app1
#NGINX Proxy prefix with a trailing slash, but proxy_pass has none
location /app1/ {
    proxy_pass http://application1;
}
#NGINX Backend prefix with a trailing slash
location /application1/ {
}

#NGINX log shows it appending the slash through too many redirects
GET /app1 HTTP/1.1 301 /app1/
GET /app1// HTTP/1.1 301 /app1//
...
#Correct response(s) should be
GET /app1/ HTTP/1.1 200 /app1/
GET /app1/ HTTP/1.1 304 /app1/
```

# Lab 2.1: Setup a Reverse Proxy

1. Create a new configuration file called `server2.conf`:

```
$ sudo vim /etc/nginx/conf.d/server2.conf
```

2. Define the following server context:

```
server {  
    listen 90;  
    root /data/server2;  
}
```

3. Save and exit the file, then open `server1.conf`:

```
$ sudo vim /etc/nginx/conf.d/server1.conf
```

# Lab 2.2: Test Your Proxy

1. Add a `proxy_pass` in `/application1` with the URI:

```
location /application1 {  
    proxy_pass http://localhost:90/sampleApp/  
}
```

2. Save and Reload NGINX

```
$ sudo nginx -s reload
```

3. Open your browser (or use `curl`) and test `http://ec2-hostname/application1`, what do you see?

```
$ curl http://ec2-hostname/application1
```

# Proxy Buffers

- `proxy_buffering on / off` sets proxy buffering
- `proxy_buffers number size` sets amount and size used for reading entire response for one connection
- `proxy_buffer_size` sets size for reading first part of a response (usually response headers) received from proxied server

## **proxy\_busy\_buffers**

This directive sets max size of “client-ready” buffers.

The “client-ready” buffers are then placed in a queue.

# Proxy Busy Buffers Example

This example increases available buffers per request, while trimming down buffers stored in responses

```
server {  
    proxy_buffering on;  
    proxy_buffer_size 1k;  
    proxy_buffers 24 4k;  
    proxy_busy_buffers_size 8k;  
    proxy_max_temp_file_size 2048m;  
    proxy_temp_file_write_size 32k;  
  
    location / {  
        proxy_pass http://example.com;  
    }  
}
```



# LOGGING

# Module Objectives

This module enables you to:

- Setup Logging to audit connections to NGINX
- Demonstrate use cases for log levels
- Differentiate between access and error log

## **error\_log** Directive

- Configures the logging settings for error messages
- Syntax: **error\_log** *file log\_level*
- "Log Level" specifies the detail of the output

# Log Levels

debug

Detailed Trace

info

General Info

notice

Something Normal

warn

Something Strange

error

Unsuccessful

crit

Important Issue(s)

alert

Fix Now!

emerg

Unusable

# **access\_log** Directive

- Records all attempts to access the server
- Default log type is *combined*

*#Example*

```
access_log /var/log/nginx/server3.access.log combined;
```

# log\_format Directive

- Defined in the `http` context
- Uses pre-defined or custom variables
- Overrides default log type i.e. `combined`

```
log_format test_log "$request $status $request_uri";

server {
    listen 80;
    root /home/student1/public_html;
    access_log /var/log/nginx/public.log test_log;
}
```

# access\_log Locations

- Default locations
  - From source: `logs/access.log`
  - Binary distro: `/var/log/nginx/access.log`

# Reading the Logs

When log values are separated by:

- , = request processed by multiple servers
- ; = internal redirect between upstreams
- 0 = unable to reach upstream
- – = internal error or cached value



# Logging Best Practices

Keep a separate log files for each server

```
server {  
    server_name server1.com;  
    root /data/server1.com;  
    error_log logs/server1.error.log    info;  
}  
  
server {  
    server_name server2.com;  
    root /data/server2.com;  
    error_log logs/server2.error.log    info;  
}
```

# Rotating Logs

Run this shell script in a *cron* job

```
#Get Yesterday's date as YYYY-MM-DD
YESTERDAY=$(date -d 'yesterday' '+%Y-%m-%d')
PID_FILE=/run/nginx.pid
LOG_FILE=/var/log/error.log
OLD_FILE=/var/log/error-$YESTERDAY.log

#Rotate yesterday's log.
mv $LOG_FILE $OLD_FILE

#Tell nginx to open the log file
kill -USR1 $(cat $PID_FILE)
```

Documentation: [Log Rotation](#)

# syslog Protocol

Use this protocol when sending messages to syslog servers such as: *splunk*, or *syslog-ng*

```
error_log syslog:server=192.168.1.1 debug;
```

```
access_log syslog:server=unix:/var/log/nginx.sock,nohostname;
```

```
access_log syslog:server=[2001:db8::1]:12345,facility=local7,tag=nginx,se
```

Documentation: [Syslog](#)

# Lab 3.1: Setup Logging

1. Open `server1.conf`

```
$ sudo vim /etc/nginx/conf.d/server1.conf
```

2. Add an `error_log` and `access_log` directive with log levels of `info` and `combined`:

```
error_log /var/log/nginx/server1.error.log info;  
access_log /var/log/nginx/server1.access.log combined;
```

3. Repeat the same steps for `server2.conf`

4. Save and Reload NGINX

```
$ sudo nginx -s reload
```

# Lab 3.2: Test Logging

1. Open a browser and test your instance:

```
$ curl http://ec2-hostname/  
$ curl http://ec2-hostname:90/
```

2. Check the logs with the `tail -f` command:

```
$ sudo tail -f /var/log/nginx/server1.access.log  
$ sudo tail -f /var/log/nginx/server2.access.log
```

# SECURITY

# Module Objectives

This module enables you to:

- Signature Strength
- Cipher Strength
- A+ Score on SSL Labs.com

# http\_ssl\_module

- Provides directives for configuring and managing HTTPS servers
- Requires `openssl` Library



# ssl vs. proxy\_ssl



# SSL Termination

SSL terminates at NGINX level, NGINX handles handshake overhead to take load off the backends

```
http {  
    ssl_session_cache    shared:SSL:10m;  
    ssl_session_timeout 10m;  
  
    server {  
        listen            443 ssl;  
        server_name       www.example.com;  
        keepalive_timeout 70;  
  
        ssl_certificate    www.example.com.crt;  
        ssl_certificate_key www.example.com.key;  
        ssl_protocols     TLSv1 TLSv1.1 TLSv1.2;  
        ssl_ciphers       HIGH:!aNULL:!MD5;  
        ...  
    }  
}
```

Documentation: [SSL Termination](#)

# Popular Use Cases

## Combine HTTP & HTTPS

```
server {  
    listen 443;  
    listen 80;  
    server_name example.com;  
    root /home/student1/public_html;  
}
```

# Popular Use Cases

Force incoming traffic to HTTPS

```
#server 1
server {
    listen 80;
    return 301 https://$host$request_uri;
}

#server 2
server {
    listen 443 ssl;
    ssl_certificate /etc/nginx/ssl/server.crt;
    ssl_certificate_key /etc/nginx/ssl/server.key;
}
```

# General Infrastructure Security

- Turn off server\_tokens
- Turn off corresponding backend engine headers
  - X-Powered-By
- Change client side error pages
- Encrypt ALL THE THINGS!!!
- Test on [SSL Labs.com](https://ssllabs.com)

# Getting a Perfect SSL Score

- Verify cert/chain in order and from trusted authority
- Use strong signature algorithm
- Use the latest protocol support
- Generate strong key signature
- Use preferred ciphers

# Configure HTTPS

- Enable `ssl` on `listen` directive
- Specify `ssl_certificate` and `ssl_certificate_key`

```
server {  
    listen 443 ssl;  
    root /data;  
  
    ssl_certificate /etc/nginx/ssl/nginx.crt;  
    ssl_certificate_key /etc/nginx/ssl/nginx.key;  
  
}
```

# ssl\_dhparam

- Specifies a file with DH parameters for DHE ciphers
- Aids in Forward Secrecy
- Make them strong, 4096 rather than 2048

```
server {  
    listen 443 ssl;  
    root /data;  
  
    ssl_certificate /etc/nginx/ssl/nginx.crt;  
    ssl_certificate_key /etc/nginx/ssl/nginx.key;  
    ssl_dhparam ssl/dhparam.pem;  
}
```



# ssl\_protocols

- Specifies which protocols are enabled
- Best Practice: Only support TLS
- Best Practice: Only support latest for higher rating

```
server {  
    listen 443 ssl;  
    root /data;  
  
    ssl_certificate /etc/nginx/ssl/nginx.crt;  
    ssl_certificate_key /etc/nginx/ssl/nginx.key;  
    ssl_dhparam ssl/dhparam.pem;  
    ssl_protocols TLSv1 TLSv1.1 TLSv1.2;  
}
```

# ssl\_ciphers

- Specifies a list based on business needs
- Moving target; ciphers constantly change, new security threats arise daily
- Configure NGINX to force client to accept preferred order of ciphers

```
server {  
    listen 443 ssl;  
    root /data;  
  
    ssl_certificate /etc/nginx/ssl/nginx.crt;  
    ssl_certificate_key /etc/nginx/ssl/nginx.key;  
    ssl_dhparam ssl/dhparam.pem;  
    ssl_protocols TLSv1 TLSv1.1 TLSv1.2;  
    ssl_ciphers AES256+EECDH:AES256+EDH:!aNULL;  
    ssl_prefer_server_ciphers on;  
}
```

# ssl\_stapling

- `ssl_stapling` "staples" an OCSP response
- `ssl_stapling_verify` verifies OCSP responses
- `ssl_trusted_certificate` required

```
server {
    listen 443 ssl;
    root /data;

    ssl_certificate /etc/nginx/ssl/nginx.crt;
    ssl_certificate_key /etc/nginx/ssl/nginx.key;
    ssl_dhparam ssl/dhparam.pem;
    ssl_protocols TLSv1 TLSv1.1 TLSv1.2;
    ssl_ciphers AES256+EECDH:AES256+EDH:!aNULL;
    ssl_prefer_server_ciphers on;
    ssl_stapling on;
    ssl_stapling_verify on;
}
```

# ssl\_session

- `ssl_session_cache`
- `ssl_session_timeout`
- `ssl_handshake_timeout`

# HSTS

- Forces browsers to communicate over **https**
- Other Security Precautions:
  - No-Content-Sniffing
  - No content displaying in iFrames

```
add_header Strict-Transport-Security max-age=63072000;  
add_header X-Content-Type-Options nosniff;  
add_header X-Frame-Options DENY;
```

# "Dual Stack" RSA and ECC

- ECC is 3x faster than RSA
- Include both pairs to support both

```
server {  
    listen 443 ssl;  
    server_name example.com;  
  
    ssl_certificate example.com.rsa.crt;    ssl_certificate_key example.c  
    ssl_certificate example.com.ecdsa.crt; ssl_certificate_key example.c  
}
```

# Elliptic Curves

NGINX doesn't have a good default for elliptic curves, so specify manually e.g.

```
ssl_ecdh_curve secp384r1;
```

# Lab 4.1: SSL Certs

1. Create a directory for certs and keys

```
$ sudo mkdir /etc/nginx/ssl -p
```

2. Change directories and run the openssl command:

```
$ cd /etc/nginx/ssl  
$ sudo openssl req -x509 -nodes -days 365 \  
-newkey rsa:2048 -keyout nginx.key -out nginx.crt \
```



# Lab 4.2: Configure SSL Parameters

1. Open `server1.conf`, and add the following directives in the `server` context:

```
ssl_certificate /etc/nginx/ssl/nginx.crt;
ssl_certificate_key /etc/nginx/ssl/nginx.key;

ssl_protocols TLSv1.2;
ssl_ciphers "AES256+EECDH:AES256+EDH:!aNULL";
ssl_prefer_server_ciphers on;
ssl_ecdh_curve secp384r1;
ssl_session_cache shared:SSL:10m;
ssl_session_timeout 10m;
ssl_session_tickets off;

add_header Strict-Transport-Security
    "max-age=63072000; includeSubdomains";
add_header X-Frame-Options DENY;
add_header X-Content-Type-Options nosniff;
```

2. Save and Close the file

# Lab 4.3: Configure HTTPS

1. Convert your `server` context to listen on `443` with `ssl`:

```
server {  
    listen 443 ssl;  
    root /home/student1/public_html;  
    error_log /var/log/nginx/server1.error.log info;  
    access_log /var/log/nginx/server1.access.log combined;  
    ...  
}
```

2. Add a new `server` block forcing `http` traffic to `https`:

```
server {  
    listen 80  
    return 301 https://$host$request_uri;  
}
```

# Lab 4.4: Test on SSL Labs.com

1. Save and **reload** NGINX
2. Test your site on **SSL Labs.com**
3. Share your results with the class



# Lab 4.5: What the "F"

1. Your test should receive an "F" rating
2. To patch the Oracle padding error, run the following commands:

```
$ sudo apt-get install --only-upgrade libssl1.0.0 openssl  
$ sudo nginx -s stop  
$ sudo apt-get update  
$ sudo nginx
```

3. Clear your cache in SSLlabs.com, and re-test
4. You should receive a "T" rating

# Restricting Access

- `allow` specifies access to a server or location block
- `deny` directive prevents access

# Basic HTTP Authentication

Requires users to specify a login and password via the browser. Syntax is:

```
auth_basic realm | off
```

# Basic Auth Example

```
server {  
    listen 8080;  
    root /home/johnny/public_html;  
  
    auth_basic "restricted server";  
    auth_basic_user_file /etc/nginx/htpasswd;  
  
    location /public_area {  
        auth_basic off;  
    }  
}
```

# Password File

- `crypt()`
- `openssl passwd`

```
name1:password1  
name2:password2  
#comments
```



# auth\_request Module

Implements authentication based on the result of a sub-request, rather than the browser

```
server {  
    listen 443 ssl;  
    root /usr/share/nginx/html;  
    ...  
    location /private {  
        auth_request /auth;  
        auth_request_set $user $upstream_http_x_user;  
        proxy_set_header x-user $user;  
        proxy_pass http://ldap_backend;  
    }  
}
```

# Lab 5.1: Create a User:Password

1. Open the password file called `htpasswd` inside `/etc/nginx`

```
$ sudo vim /etc/nginx/htpasswd
```

2. Delete the password for the admin entry
3. Close and save the file. Back in the terminal, run the following command:

```
$ openssl passwd
```

4. Specify any user and password, copy the encrypted output to your clipboard, and paste it back into `htpasswd`

# Lab 5.2: Setup Basic Auth

1. Open `server1.conf` and enable authentication on the `server` context. Then point to the `htpasswd` file:

```
$ sudo vim /etc/nginx/conf.d/server1.conf

server {
    listen 443 ssl;
    ...
    auth_basic "protected";
    auth_basic_user_file /etc/nginx/htpasswd;
}
```

2. Save and reload NGINX. Open a browser and test your instance's home page
3. Login with your username and password you specified

# Setup Limit Rates

Limitations based on:

- Number of Connections
- Number of Requests
- Download Speeds

# Limiting Connections

- Define a Shared Memory Zone
- `limit_conn_zone` defines: zone, size, name, and key
- Syntax is:

```
limit_conn_zone $variable zone=name:size;
```

# limit\_conn

Specify shared memory zone in desired context

Syntax:

```
limit_conn zone number_of_connections;
```

```
http {  
    limit_conn_zone $binary_remote_addr zone=ip:10m;  
  
    server {  
        ...  
        location /downloads {  
            limit_conn ip 5;  
        }  
    }  
}
```

# Limiting Request Rate

Set request rate, name, key, and duration in Shared Memory Zone

Example:

```
limit_req_zone $server_name zone=one:10m rate=1r/s
```

# limit\_req

Syntax:

```
limit_req zone='name' [burst=n.o | nodelay];
```

```
http {  
    limit_req_zone $server_name zone=ten:10m rate=10r/s  
    ...  
    server {  
        ...  
        location /data {  
            limit_req zone=ten burst=15;  
        }  
    }  
}
```



# Limiting Download Rate

Syntax:

*limit\_rate speed*

```
location /media {  
    limit_conn ipzone 5;  
    limit_rate 50k;  
}
```

# limit\_rate\_after

Throttle download speeds after client builds a buffer

```
location /videos {  
    limit_rate_after 500k;  
    limit_rate 50k;  
}
```

# VARIABLES

# Module Objectives

This module enables you to:

- Gain knowledge of NGINX's predefined variables
- Understand Variable Scope with regards to NGINX
- Define your own custom variables

# Core Module Variables

Variable	Value
<code>\$host</code>	Host name from request line
<code>\$request_uri</code>	The full URI, including arguments
<code>\$uri</code>	Normalized URI (no arguments)
<code>\$scheme</code>	Request scheme (HTTP or HTTPS)
<code>\$request_method</code>	GET, POST, PUT etc.
<code>\$request_filename</code>	Absolute file path for current request
<code>\$remote_addr</code>	IP address of client

# Variable Example

Given URL: `http://localhost:8080/test?arg=1`

- `$host` = `localhost:8080`
- `$request_uri` = `/test?arg=1`
- `$uri` = `/test`
- `$scheme` = `http://`
- `$args` = `?arg=1`

Usage Example:

```
server {  
    listen 80;  
    return 301 https://$host$request_uri;  
}
```

# Useful Troubleshooting Variables

- `$upstream_connect_time`
- `$upstream_header_time`
- `$upstream_response_time`
- `$request_time`

# set Directive

Syntax:

`set $variable_name value;`

```
set $foo hello;  
set $bar "hello world";  
set $combo $foo
```



# Declaration and Scope

- Declared when NGINX loads the configuration file
- Assigned at runtime
- Every config can have access to a variable
- Similar to Java method invocation

# Scope Example

```
server {  
    listen 8080;  
  
    location /test1 {  
        return 200 = "foo $foo \name example = $example \n";  
    }  
  
    location /test2 {  
        set $foo hello;  
        return 200 = "foo = $foo \n";  
    }  
}  
  
server {  
    listen 8081;  
  
    set $example 42;
```

# map Directive

Creates a mapping relationship between two variables

Syntax:

```
map $var1 $var2 {value value}
```

```
map $uri      $path {
    /test1    /path1
    /test2    /path2
    /test3    /path3
}

server {
    listen 80;

    location /test1 {
        return 200 "$path \n";
    }
}
```

# Default Value

Default values protect NGINX from translation errors

```
map $arg $value {  
    default 1;  
    test1 2;  
}
```

# Regex in Maps

Only named captures work in maps, no positional captures

```
map $uri $path {  
    ~*/test/.*\.php$ /path4;  
    /example /examplePath;  
}
```

# Use Case: Conditional Logging

Use variables to determine "loggability"

```
map $status $loggable {  
    ~^[23] 0;  
    default 1;  
}  
  
access_log /path/to/access.log combined if=$loggable;
```

# Lab 6: Map Example

1. Create a **map** for variables **\$is\_redirect** and **\$uri**:

```
map $uri $is_redirect {  
    default      0;  
    /test1      1;  
    /test2      2;  
    /test3      3;  
}
```

2. Create a regex for /test:

```
server {  
    ....  
    location ~* /test(\d+)$ {  
        return 200 "variable = $is_redirect \n";  
    }  
}
```

3. Test the map using curl. For example:

```
curl http://localhost/test1
```

**END OF PART ONE:**

**THANK YOU!**



# Agenda: Part Two



- Routing Connections
- Load Balancing
- Caching
- Compression
- Dynamic Configuration
- Installation

# ROUTING CONNECTIONS

# Module Objectives

This module enables you to:

- Use specific directives in NGINX to reroute traffic
- Define URL Rewrites
- Understand Rewrite Request processing

# alias Directive

Allows you to specify a replacement path

Syntax:

**alias** *path*;

```
server {  
    root /home/public_html;  
  
    location /test {  
        alias /data/app1;  
    }  
}
```

# Lab 7: Alias Replacement

1. Open `server1.conf`:

```
$ sudo vim /etc/nginx/conf.d/server1.conf
```

2. In the `/application2` prefix, specify a replacement path of `/data/test` using the `alias` directive:

```
location /application2 {  
    alias /data/test;  
}
```

3. Save the file and reload Nginx. Test against `http://server/application2/logo.png` What do you observe? Try the other URIs

# alias with Regex

```
location ~ ^/pictures/(.+\.(:gif|jpe?g|png))$ {  
    alias /data/images/$1;  
}
```

# Lab 7.5: Regex Alias

1. Open `server1.conf` and edit the `/images` prefix to be a case insensitive regex with a URI match of `/pictures`:

```
location ~ ^/pictures/(.+\.?(gif|jpe?g|png))$ {  
    alias /data/images/$1;  
}
```

2. Save the file and reload Nginx. Test against `http://server/pictures/logo.png` What do you observe?

# return Directive

Return a HTTP response code and URL to the client

Syntax:

*return code url;* Or *return url;*

```
server {  
    listen 8080;  
    root /home/public_html;  
  
    location /test {  
        return http://localhost:8081$uri;  
    }  
  
server {  
    listen 8081;  
    root /data/app1;  
    autoindex on;  
    }  
}
```



# rewrite Directive

Regex pattern matched against URI, replacement string rewrites the URI.

Syntax:

*rewrite regex replacement [flag]*

```
server {  
    listen 8080;  
    root /home/public_html;  
    rewrite ^/shop/products/(\d+) /myshop/products/product$1.html;  
}
```

# Lab 8.1: Setting Up Rewrite Data

1. Change directories to:

```
$ cd /home/student1/public_html/shop/products
```

2. Ensure that the following files exist:

```
$ ls  
product1.html product2.html product3.html
```

3. Open `product1.html` and edit the paragraph tag by removing `8080` and replacing `<server>` with your ec2-hostname

```
<p>  
  
</p>
```

# Lab 8.2: Rewrite URLs

1. Open `server1.conf` in `/etc/nginx/conf.d`
2. Inside the `server` context, define a `rewrite` regex:

```
^/shop/greatproducts/(\d+)$
```

and a replacement string:

```
/shop/products/product$1.html
```

3. Save your file and reload nginx
4. Open your browser and test against:
  - `<server>/shop/greatproducts/2`
  - `<server>/shop/greatproducts/3`
  - `<server>/shop/greatproducts/1`

# Lab 8.3: Rewrite URLs (Continued)

1. Now try `<server>/shop/greatproducts/1`
2. What do you notice about the image?
3. Add another `rewrite` at the `server` level with the following regex:

```
^/media/pics/(.+\. (gif|jpe?g|png))$
```

and replacement string:

```
/images/$1
```

4. Save and reload NGINX
5. Re-test `/shop/greatproducts/1`
6. If the image still doesn't render, ensure you have a prefix or regex location enabled for your image content

# Rewrite Process Cycle

## Highlights:

- Executed sequentially
- Executed upon server selection
- All rewrites in higher context are executed first
- Flags will stop further processing

# rewrite flags

Prevents further rewrite processing

- `last`
- `break`
- `redirect`
- `permanent`

# Lab 9.1: Rewrite Flags

1. Open `server1.conf` and define a new location block with the prefix `/shop`
2. Cut and paste the `rewrite` regarding `/greatproducts`, and paste it into the new `location`
3. Add a new `rewrite` with regex:

```
^/shop/.+/(\\d+)$
```

and replacement string:

```
/shop/services/service$1.html
```

4. Add a `return` directive with code `403`
5. Save and reload NGINX. Re-test  
`/shop/greatproducts/1`

## Lab 9.2: Rewrite Flags (Continued)

1. Re-open `server1.conf` and place `break` flags after every `rewrite` in the `location` context
2. Save and reload NGINX
3. Re-test `/shop/greatproducts/1`



# LOAD BALANCING

# Module Objectives

This module enables you to:

- Setup basic load balancer
- Identify load balancing methods
- Enable session persistence

# ngx\_http\_upstream\_module

Key Take-Aways:

- Defines a group of servers
- Leverages `proxy_pass` directive
- Server definition can be:
  - Unix socket
  - DNS
  - IP:Port

# Load Balancing

`proxy_pass` forwards request to `upstream` link

```
upstream myServers {  
    server training.example.com;  
    server training.example1.com:8080;  
    server 192.168.245.27;  
}
```

# Specifying Server Priorities

`weight` indicates `server` priority

`max_fails` indicates server-level failures

`fail_timeout` indicates timeout and duration of downtime

```
upstream myServers {  
    server backend.server1 weight=5 max_fails=10 fail_timeout=90s;  
    server backend.server2 weight=3 max_fails=4 fail_timeout=60s;  
    server backend.server3 weight=4 max_fails=2 fail_timeout=30s;  
}
```

# Lab 10.1: Setup "Backends"

1. In the `/etc/nginx/conf.d`, create a new `.conf` called `backends.conf` and define a `server` context with the `root` directive pointing to `/data/backend1`. The server should listen on `8081`
2. Repeat steps 1 and 2 for the remaining servers. Servers should listen on `8082` and `8083` respectively, and the `root` directories are `/data/backend2` and `/data/backend3`

# Lab 10.2: Configure Load Balancing

1. Create a file called `myServers.conf` and define an `upstream` with three backend servers using `127.0.0.1:<port>`
2. Create a `server` context that listens on `8080`
3. Set the `root` directive to your `public_html` directory.
4. Define an `error_log` with a level of `info` and an `access_log` with a level of `combined`
5. Create a `location` context that matches all requests
6. Add a `proxy_pass` to forward all request to the `upstream`

# ngx\_stream\_core\_module

## Key Take-Aways:

- Used for TCP/UDP Load Balancing
- Also leverages `proxy_pass` directive
- Similar syntax with `ngx_http_upstream` module
- Version compatibility:
  - TCP: r5 or greater
  - UDP: r9 or greater
- Exists in `main context`



# **stream** Use Cases

- **TCP:** mSQL, LDAP, RTMP
- **UDP:** DNS, Syslog, RADIUS

# mySQL Example

Load Balancings across three SQL servers

```
stream {  
    server {  
        listen 3306;  
        proxy_pass db;  
    }  
  
    upstream db {  
        server db1:3306;  
        server db2:3306;  
        server db3:3306;  
    }  
}
```

# DNS Example

Load balances UDP traffic across two DNS servers

```
stream {
    upstream dns_upstreams {
        server 192.168.136.130:53;
        server 192.168.136.131:53;
    }

    server {
        listen 53 udp;
        proxy_pass dns_upstreams;
        proxy_timeout 1s;
        proxy_responses 1;
        error_log logs/dns.log;
    }
}
```

# Load Balancing Methods

- `least_conn`
- `least_time`
- `hash`
- `ip_hash`

# least\_conn Directive

```
upstream backendServers {  
    least_conn;  
  
    server backend1.com;  
    server backend2.com;  
    server backend3.com;  
}
```

# least\_time Directive

```
upstream myServers {  
    least_time header;  
  
    server backend1.com;  
    server backend2.com;  
    server backend3.com;  
}
```

# hash Directive

```
upstream myServers {  
    hash $request_uri;  
  
    server backend1.com;  
    server backend2.com;  
    server backend3.com;  
}
```

# **ip\_hash** Directive

Uses first three octets for IPv4, or entire IPv6 address

```
upstream myapp1 {  
    ip_hash;  
    server srv1.example.com;  
    server srv2.example.com;  
    server srv3.example.com;  
}
```



# fail\_timeout Parameter

```
upstream myServers{  
    server backend1.example.com:8080 max_fails=3 fail_timeout=30s;  
    ...  
}
```

# max\_conns Parameter

```
upstream backend {  
    server backend1.example.com max_conns=3;  
    server backend2.example.com;  
}
```

# queue Directive

```
upstream backend {  
    server backend1.example.com max_conns=3;  
    server backend2.example.com;  
        queue 100 timeout=70;  
}
```

# Session Affinity

For applications that require state data on backend servers

NGINX supports the following methods:

- `sticky cookie`
- `sticky learn`
- `sticky route`

# HTTP using DNS

NGINX Plus can monitor changes of IP addresses that correspond to domains and resolve them

```
http {
    resolver 10.0.0.1 valid=300s ipv6=off;
    resolver_timeout 10s;

    server {
        location / {
            proxy_pass http://backend;
        }
    }

    upstream backend {
        zone backend 32k;
        least_conn;
        ...
        server backend1.example.com resolve;
        server backend2.example.com resolve;
    }
}
```

# Load Balancing NTLM

Release 7+ only. Requires HTTP 1.1 connection

```
http {
    ...
    upstream exchange {
        zone exchange 64k;
        ntlm;
        server exchange1.example.com;
        server exchange2.example.com;
    }

    server {
        listen          443 ssl;
        ssl_certificate  /etc/nginx/ssl/company.com.crt;
        ssl_certificate_key /etc/nginx/ssl/company.com.key;
        ssl_protocols   TLSv1 TLSv1.1 TLSv1.2;

        location / {
            proxy_pass      https://exchange;
            proxy_http_version 1.1;
        }
    }
}
```

# Lab 11: Configure Hash Method

1. Open `myServers.conf` and define a `hash` selection algorithm in `upstream myServers`
2. Save and reload nginx, test in browser several times
3. Notice how the session sticks to that particular server?
4. Now run a dynamic curl request from a local terminal

```
for i in `seq 1 100` ; do curl -s -o /dev/null -w "%{http_code}" \
http://<external ip>/\?$i \
; done
\
```

5. Check the log files; are the connections still persistent?

# **LIVE ACTIVITY MONITORING**



# Module Objectives

This module enables you to:

- Use the `status` directive to get server metrics
- Configure `health_check` to monitor the availability of your upstream servers

# status Directive

```
server {  
    listen 8080;  
  
    location = /status{  
        status;  
    }  
}
```

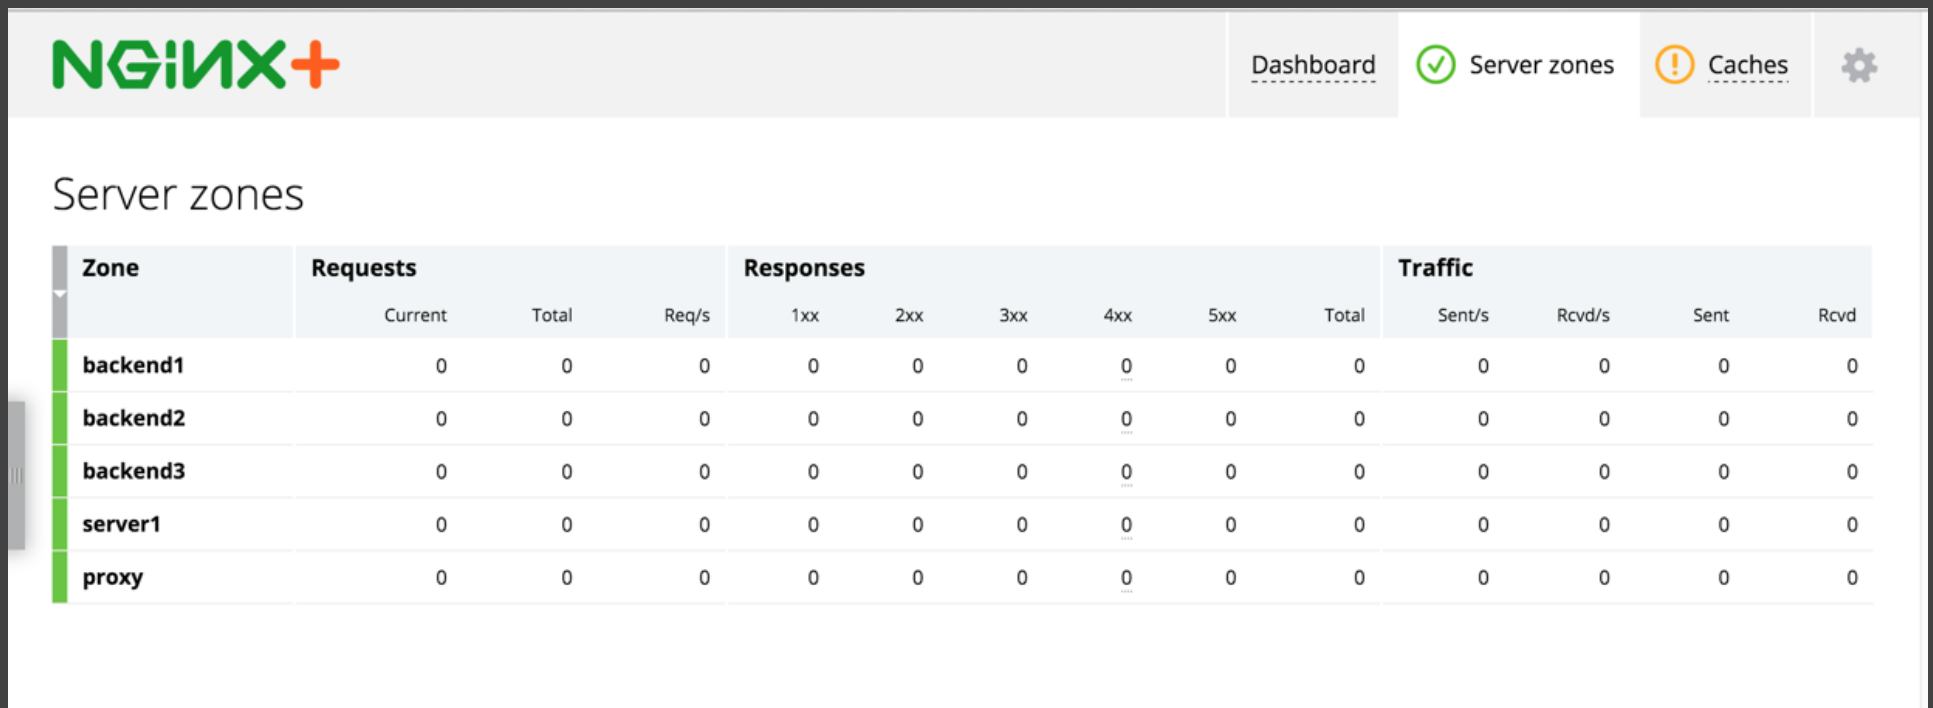
# Securing the Status Request

```
location = /status{  
    allow 192.168.0.0/16;  
    deny all; # deny access from everywhere else  
  
    status;  
}
```

# Default Status Page

NGINX Plus contains an html page that parses JSON

The page is located at **`/usr/share/nginx/html/status.html`**



The screenshot shows the NGINX Plus status page. At the top is the NGINX+ logo. To the right are navigation links: 'Dashboard' (underlined), 'Server zones' (with a green checkmark icon), 'Caches' (with a yellow warning icon), and a settings gear icon. Below the navigation bar, the title 'Server zones' is displayed. A table follows, showing metrics for five zones: backend1, backend2, backend3, server1, and proxy. The table has four main sections: Zone, Requests, Responses, and Traffic. Each section contains specific metrics as detailed in the table below.

Zone	Requests			Responses						Traffic			
	Current	Total	Req/s	1xx	2xx	3xx	4xx	5xx	Total	Sent/s	Rcvd/s	Sent	Rcvd
backend1	0	0	0	0	0	0	0	0	0	0	0	0	0
backend2	0	0	0	0	0	0	0	0	0	0	0	0	0
backend3	0	0	0	0	0	0	0	0	0	0	0	0	0
server1	0	0	0	0	0	0	0	0	0	0	0	0	0
proxy	0	0	0	0	0	0	0	0	0	0	0	0	0

# Server Zones

```
server {  
    listen 8081;  
    root /server/backend1;  
    status_zone apac;  
}
```

```
server {  
    listen 8082;  
    root /server/backend2;  
    status_zone apac;  
}
```

```
server {  
    listen 8083;  
    root /server/backend3;  
    status_zone us;  
}
```

Individual servers zones displayed in status.html:

Server zones

Zone	Requests			Responses			
	Current	Total	Req/s	1xx	2xx	3xx	4xx
apac	0	0	0	0	0	0	0
us	0	0	0	0	0	0	0

# zone Directive

Shared memory zone makes upstream dynamically configurable

Allows worker processes to share counter information

```
upstream myServers {  
    zone backend 64k;  
    server backend1;  
    server backend2;  
    server backend3;  
}
```

# Lab 12.1: Define a Status Page

1. Open `myServers.conf` located in `/etc/nginx/conf.d/`
2. In a new `server` context that listens on `9090` with a `root` location `/usr/share/nginx/html;`
3. Add a `location` prefix with an exact match `= /status`, and the `status` directive in this location block
4. Save and reload NGINX
5. Open a browser and access the following URI:

```
http://<server>:9090/status.html
```

6. Notice anything strange in the GUI?

# Lab 12.2: Define Server Zones

1. In `backends.conf`, specify a server zone for each server by using the `status_zone` directive
2. Open `server1.conf` and define a `status_zone`
3. Open `server2.conf` and define a `status_zone`
4. Open `myServers.conf` and specify a shared memory zone in your `upstream` block with a size of `64k`
5. Save and reload NGINX
6. Refresh your browser listening on

```
http://<server>:9090/status.html
```



# Server health\_check

A request sent to upstream to check status based on conditions

```
upstream myUpstreams {
    zone backend 64k;

    server localhost:8081 weight=1;
    server localhost:8082 weight=4;
    server localhost:8083 weight=7;
}

server {
    listen 8080;
    error_log /var/log/nginx/upstream_error.log info;

    location / {
        proxy_pass http://myUpstreams;
        health_check;
    }
}
```

# health\_check Parameters

- interval
- fails
- passes
- uri
- match

# match Block

Block directive that defines conditions for:

`health_check`

Conditions can be based on:

- Response Codes
- Header Values
- Text body of documents

# match Directives

- status
  - status 200
  - status ! 403
  - status 200-399
- header
  - header Content-Type = text/html
  - header Cache-Control
- body
  - body ~ "Hellow World"
  - body !~ "hello world"

# match Example

```
server {  
    location / {  
        proxy_pass http://myServers;  
        health_check match=conditions fails=2;  
    }  
}  
  
match conditions {  
    status 200;  
    header Content-Type = text/html;  
    body !~ "maintenance";  
}
```

# Lab 13.1: Server Maintenance

1. Open the `myServers.conf`, file
2. Change your upstream `server` entries back to `localhost:8081, localhost:8082, localhost:8083`
3. Define a `match` block called `health_conditions` with the following directives:

```
match health_conditions {  
    status 200-399;  
    header Content-Type = text/html;  
    body !~ maintenance;  
}
```

# Lab 13.2: Health Check Params

1. In the `/`, add the following `match` parameters:

```
location / {  
    proxy_pass http://myServers;  
    health_check match=health_conditions  
    fails=2  
    uri=/health/test.html;  
}
```

2. Save and reload NGINX
3. Refresh your browser listening on

```
http://<server>:9090/status.html
```

4. A server is down. To fix it, change the text "maintenance" in the test.html for backend3

```
$ cd /data/server2/backend3/health/  
$ sudo vim test.html
```

# CACHING



# Module Objectives

This module enables you to:

- Define a reverse proxy cache for your upstream and other servers
- Purge old or stale content from the cache
- Identify other cache control techniques

# Reverse Proxy and Caching

Common use case to have NGINX in front, caching static resources to improve performance

To compose a cache:

- Define a `cache_path`
- Configure the `proxy_pass`
- Reference the `cache_key`
- Validate the cacheability of content using `proxy_cache_valid`

# proxy\_cache\_path

- Defined in `http` context
- Directive must point to directory or mount point/volume attached to instance
- Uses `hybrid persistent` model via a two level directory system

```
proxy_cache_path /data/nginx/cache levels=1:2 keys_zone=img_cache:20m in:
```

# Configuring the **proxy\_cache**

- **proxy\_cache\_key**: generates md5 hash
- **proxy\_cache**: writes to the cache

```
location / {  
    proxy_pass http://application.com:8080;  
  
    proxy_cache_key "$scheme$host$request_uri";  
    proxy_cache my-cache;  
}
```

# Validating the Cache

Honors (or overrides) Cache-Control origin headers

```
location / {  
    proxy_pass http://application.com:8080;  
  
    proxy_cache_key "$scheme$host$request_uri";  
    proxy_cache my-cache;  
  
    proxy_cache_valid any 1m;  
    proxy_cache_valid 404 1m;  
}
```

# Passing Headers

In order to audit client address, as well as proxy address, we can forward the headers to the backend

```
proxy_set_header    Host      $host;  
proxy_set_header    X-Real-IP  $remote_addr;  
proxy_set_header    X-Forwarded-For $proxy_add_x_forwarded_for;
```

# add\_header Directive

Inserts response headers—can map to a variable

```
server {  
  [...]  
    add_header X-Proxy-Cache $upstream_cache_status;  
  
    location / {  
      proxy_cache myCache;  
      proxy_pass http://localhost:8081;  
    }  
}
```

# Lab 14.1: Proxy Cache

1. Open `server1.conf`
2. Define a cache path in the `http` context:

```
proxy_cache_path /data/nginx/cache levels=1:2  
keys_zone=img_cache:20m inactive=5m;
```

3. In the `server` context, set the `proxy_cache_key` to the `$scheme`, `$host`, and `$request_uri`
4. Add the following in your `server` context

```
proxy_cache_key $scheme$host$request_uri;  
proxy_set_header Host $host;  
proxy_set_header X-Real-IP $remote_addr;  
proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
```



# Lab 14.2: Proxy Cache Continued

1. Set the `proxy_cache` in the `application1` prefix, and set the validation of the cache for 10 minutes.

```
location /application1 {  
    proxy_cache img_cache;  
    proxy_cache_valid 10m;  
    proxy_pass http://localhost:90/sampleApp
```

2. Save and reload NGINX
3. Reload NGINX and make a request to

```
http://<server>/application1
```

4. Hit refresh multiple times and then check your `status.html` in a separate tab. Notice the cache icon is now “warm” and the hit ratio is increasing

# Lab 15.1: Proxy Upstream Cache

1. Open `Open /etc/nginx/conf.d/myServers.conf`
2. Create a `proxy_cache_path` with a `keys_zone` named `upstreamCache` that lasts for:
  - 10 minutes
  - Has a max size of 60mb
  - Timeouts after 60 minutes

# Lab 15.2: Proxy Upstream Cache Continued

1. Set the `proxy_set_header` to forward the client host and IP
2. Use the `add_header` directive with the `X-Proxy-Cache` response header to return the upstream status
3. Add `proxy_cache` to the proxying location
4. Save and reload NGINX
5. Make a request to your upstream and notice the second cache in your status.html
6. Try making a `curl` request? What do you see with the `-I` parameter?

# Caching Resources

Directives that control cached responses:

- `proxy_cache_min_uses`
- `proxy_cache_methods`

Caching limit rates:

- `proxy_cache_bypass`
- `proxy_no_cache`

Documentation: [Cache Admin Guide](#)

# proxy\_cache\_min\_uses

```
server {  
    proxy_cache myCache;  
    proxy_pass http://localhost:8081;  
    proxy_cache_min_uses 5;  
}
```

# cache\_methods and no\_cache

Syntax:

`proxy_cache_methods <REQUEST METHOD>`

Syntax:

`proxy_no_cache $arg`

```
map $request_uri $no_cache;  
    /default      0;  
    /test         1;  
  
server {  
    proxy_cache_methods GET HEAD POST;  
    proxy_no_cache $no_cache;  
}
```

# Cache Manager and Loader

- `loader_threshold`
- `loader_files`
- `loader_sleeps`

```
proxy_cache_path /data/nginx/cache keys_zone=one:10m loader_threshold=300
```

# proxy\_cache\_purge Directive

Allows you to remove full cache entries that match a configured value.

```
server {  
    proxy_cache myCache;  
    proxy_pass http://localhost:8081;  
    proxy_cache_purge $purge_method;  
}
```



# Purge Methods

## Partial Purge

- use `curl` command to send `PURGE HTTP` request, `map` evaluates request and enables the directive

## Full Purge

- turn `purger` parameter on in the `proxy_cache_path`, all wildcard pages will also be purged

# HTTP PURGE Example

Request:

```
$curl -X PURGE -D - "http://www.mysite.com"
```

```
# setting the default purge method will only delete matching URLs.
map $request_method $purge_method {
    PURGE 1;
    default 0;
}
server {
    listen 80;
    server_name www.mysite.com
    proxy_cache myCache;
    proxy_pass http://localhost:8081;
    proxy_cache_purge $purge_method;
}
```

# purger Example

Request:

```
$curl -X PURGE -D - "http://www.mysite.com/*"
```

```
proxy_cache_path /data/nginx/cache levels=1:2 keys=myCache:10m purger=on
```

```
server {  
    listen 80;  
    server_name www.mysite.com;  
    location / {  
  
        proxy_cache_purge $purge_method;  
    }  
}
```

# Lab 16.1: Configure Cache Purge

1. Open `myServers.conf` and use a `map` to create a custom variable called `purge_method` that depends on the predefined `request_method`
2. In the `location` where caching occurs, specify a condition for the cache purge request.
3. Save and reload NGINX
4. Send the purge command using the `curl` command, your machine url, and port. E.g. `http://<server>:8080.`
5. A successful purge should return an `HTTP 204` code (no content).

# COMPRESSION

# HTTP **gzip** Module

Key Directives:

- gzip
- gzip\_types
- gzip\_proxied

# gzip Example

```
http {  
    gzip on;  
    gzip_types text/plain text/css;  
    gzip_proxied any;  
}
```

# gzip\_min\_length

Specifies the minimum length of the response to compress

```
gzip_min_length 1000;
```



# gzip\_proxied

NGINX doesn't compress proxied requests by default,

**gzip\_proxied** instructs NGINX to check header fields

```
server {  
    gzip on;  
    gzip_types text/plain application/xml;  
    gzip_proxied no-cache no-store private expired auth;  
    gzip_min_length 1000;  
}
```

# gzip\_static

Syntax:

`gzip_static on | off`

```
server {  
  gzip on;  
  gzip_static on;  
  gzip_types text/plain application/xml;  
  gzip_proxied no-cache no-store private expired auth;  
  gzip_min_length 1000;  
}
```

# HTTP **gunzip** Module

Decompresses client gzip responses, if gzip method isn't supported

Syntax:

**gunzip\_buffers** *number size*

```
http {  
    gunzip on;  
    gunzip_buffers 32 4k;  
}
```

# gzip\_vary

Places the

“Vary: Accept-Encoding” response header if both `gzip_static` and `gunzip` are active.

```
server {  
    gzip on;  
    gzip_vary on;  
    gzip_types text/plain application/xml application/json;  
    gzip_proxied no-cache no-store private expired auth;  
    gzip_min_length 1000;  
}
```

# **DYNAMIC CONFIGURATION**

# Module Objectives

This module enables you to:

- Leverage NGINX API to dynamically configure server information
- Use the state directive to make changes persistent

# Dynamic Configuration

## Advantages:

- View, modify, and remove servers at runtime
- No need to reload NGINX to affect changes

## Disadvantages

- Runtime config is not saved to conf file
- Changes revert back to conf settings after reload (unless using `state` directive)

# Shared Memory Zone

- Required to make server dynamically configurable
- Distributes traffic more evenly
- Necessary for `state` changes (covered later)



# upstream\_conf Directive

```
upstream myServers {  
  
    server localhost:8081;  
    server localhost:8082;  
}  
  
server {  
    listen 8080;  
  
    location / {  
        proxy_pass http://myServers;  
    }  
  
    location /upstream_conf {  
        upstream_conf;  
        allow 127.0.0.1;  
        deny all;  
    }  
}
```

# `upstream_conf` Parameters

Key Parameters:

- `add=`
- `remove=`
- `drain=`
- `upstream=name`
- `server=address`
- `id=number`

# Sending Requests

View and modify server details

```
#View all primary servers in upstream group myServers  
curl http://<server>:8080/upstream_conf?upstream=myServers
```

```
#View individual server detail in upstream group  
curl http://<server>:8080/upstream_conf?upstream=myServers&id=<id number>
```

# Add and Remove Servers

```
#Add a new server to the myServer group with address localhost:8083 and \
http://<server>:8080/upstream_conf?add=&upstream=myServers&server=localhost
```

```
#Remove server with id=0 from the myServers upstream group
http://<server>:8080/upstream_conf?remove=&upstream=myServers&id=0
```


# More Examples

#Modify the server with id = 0 and set the weight to 5 and the max\_fails  
`http://<server>:8080/upstream_conf?upstream=myServers&id=0&weight=5&max_:`

#Modify the server with id = 0 and set the route parameter to tomcat1  
`http://<server>:8080/upstream_conf?upstream=myServers&id=0&route=tomcat1`


#Modify the server with id = 0 and set the server address to newdomain.co  
`http://<server>:8080/upstream_conf?upstream=myServers&id=0&server=newdoma`

# Using the Status Dashboard

http\_backends.test 

Edit selected

Add server

	Server			Requests		Responses			
	<input type="checkbox"/>	Name	DT	W	Total	Req/s	...	4xx	5xx
	<input type="checkbox"/>	95.211.80.227:80	1d 1h	1	0	0		0	0
	<input type="checkbox"/>	95.211.80.227:80	1d 1h	1	0	0		0	0
	<input type="checkbox"/>	127.0.0.1:8080	1h 18m	1	0	0		0	0

# Lab 17: Dynamic Config

1. Open `myServers.conf` and add an `upstream_conf` location prefix in the same server block where the `status` prefix is located.
2. Add the `upstream_conf` directive inside the `upstream_conf` prefix
3. Open your browser and hit `http://server:9090/upstream_conf?upstream=myServers`, note the id number of the first server
4. Use the API commands to `remove` the server, then `add` it back with a `weight` of 5
5. Try using the `status.html` page to change server details

# Persistent Changes

For SDP protocols that allow automatic detection of devices and/or services, changes must persist across reloads

Documentation: [Microservices architecture](#)

Demo: [Service Discovery with Consul](#)



# state Directive

**state** name MUST match **zone** name

Syntax:

**state** *file/path.state*

```
upstream myServers {  
    zone backend 64k;  
    state /etc/nginx/conf.d/backend.state;  
}
```

# Lab 18.1: Persistent Changes

1. Create a directory for the `state` file and change ownership so NGINX has write permissions:

```
$ sudo mkdir -p /var/lib/nginx/state  
$ sudo chown nginx:nginx /var/lib/nginx/state
```

2. Open `server1.conf` and comment out each `server` in the `upstream`, then add the `state` directive:

```
upstream myServers {  
    zone backend 64k;  
    state /var/lib/nginx/state/backend.state;  
  
    #server 127.0.0.1:8081;  
    #server 127.0.0.1:8082;  
    #server 127.0.0.1:8083;  
}
```

# Lab 18.2: Changing the State

1. Save NGINX and reload. Check `status.html`
2. Your servers vanished!
3. Create/edit the state file and add server details to create a new upstream:

```
#Create and edit the state file:  
$ sudo vim /var/lib/nginx/state/backend.state  
  
server 127.0.0.1:8081;  
server 127.0.0.1:8082;  
server 127.0.0.1:8083;
```

4. Go to your `status.html` page and edit these server details using the GUI
5. Reload and NGINX and see if the changes persist

# INSTALLATION

# Module Objective

This module will enable you to:

- Install NGINX from a binary distribution
- Compile binary from source code

# CentOS, RHEL

- Setup **yum** repository
- Edit repo file to pull latest packages
- Update repo
- Install NGINX

Documentation: [CentOS/RHEL Install Guide](#)

# Debian, Ubuntu

- Authenticate repo signature
- Retrieve distribution components
- Resolve dependencies
- Install NGINX

# NGINX Signing Key

```
wget http://nginx.org/keys/nginx_signing.key  
sudo apt-key add nginx_signing.key
```



# Distribution URL

Edit the sources.list to retrieve correct distribution components

```
#Open the sources.list file with vim  
sudo vim /etc/apt/sources.list
```

```
#For Debian, append the following distribution URLs  
deb http://nginx.org/packages/debian/ codename nginx  
deb-src http://nginx.org/packages/debian/ codename nginx
```

```
#For Ubuntu  
deb http://nginx.org/packages/ubuntu/ codename nginx  
deb-src http://nginx.org/packages/ubuntu/ codename nginx
```

# Distribution Codename Reference

Debian:

- 7.x - wheezy
- 6.x - squeeze

Ubuntu:

- 16.04 - xenial
- 14.04 - trusty

```
deb-src http://security.ubuntu.com/ubuntu xenial-security main restricted
deb http://security.ubuntu.com/ubuntu xenial-security universe
deb-src http://security.ubuntu.com/ubuntu xenial-security universe
deb http://security.ubuntu.com/ubuntu xenial-security multiverse
deb-src http://security.ubuntu.com/ubuntu xenial-security multiverse

## Uncomment the following two lines to add software from Canonical's
## 'partner' repository.
## This software is not part of Ubuntu, but is offered by Canonical and the
## respective vendors as a service to Ubuntu users.
# deb http://archive.canonical.com/ubuntu xenial partner
# deb-src http://archive.canonical.com/ubuntu xenial partner
deb http://nginx.org/packages/ubuntu/ xenial nginx
deb-src http://nginx.org/packages/ubuntu/ xenial nginx
```

# Install: apt-get

```
sudo apt-get update  
sudo apt-get install nginx
```

# Check Installation

NGINX will run on port 80 by default



# Location of Files

## NGINX Executable

```
/usr/sbin/nginx
```

## Configuration File

```
/etc/nginx
```

## Log Files

```
/var/log/nginx
```

# Building NGINX From Source

## General Steps:

- Download `.tar` file
- Extract the archive
- Run the `.configure` tool
- Add modules with various parameters using:

```
./configure --<param>=<paramValue>
```

- Run `make && sudo make install`

# Command Steps

Copy the Mainline download link here:

<http://nginx.org/en/download.html>

```
$ sudo wget http://<nginx_mainline_verison>
$ tar -xvf <nginx_mainline_version>.tar.gz
$ sudo apt-get install libpcre3-dev build-essential libssl-dev
$ cd <nginx_mainline_verison>
$ ./configure --with-http_ssl_module --with-debug --with-<other_modules>
$ make && sudo make install
```

# Important Notes

Make sure to download requisite libraries PRIOR to compiling the binary

Specify file paths as needed:

- `--prefix=path`
- `--sbin-path=path`
- `--conf-path=path`
- `--error-log=path`
- `--http-log=path`

```
# Example  
./configure --sbin-path=/usr/local/nginx --error-log=/logs
```



# **ADDITIONAL RESOURCES**

# Further Information

- [NGINX Documentation](#)
- [NGINX Admin Guides](#)
- [NGINX Blog](#)

# Q&A

- Survey!
- Sales: [nginx-inquiries@nginx.com](mailto:nginx-inquiries@nginx.com)